

LLM Synthesis and Citation Grounding

Prompt Engineering for Synthesis

Prompt templates instruct the LLM to synthesize concise answers directly referencing bracketed citations [Doc X, Chunk Y] corresponding to retrieved context.

Streaming Responses & Server-Sent Events

RAG systems stream response tokens back to users via Server-Sent Events (SSE) or WebSockets, reducing Time-To-First-Token (TTFT) to under 300ms.

Hallucination Mitigation

Post-processing verification checks generated claims against source citations, redacting ungrounded statements prior to client delivery.